

Machine Learning 2: Maternal Health Risk Classification

Aisha, Mufaddal, Raju Ahmed

2026-01-02

Abstract

Maternal health is a major challenge in Bangladesh, particularly in rural areas where healthcare is hard to access. Many pregnant women suffer from conditions like high blood pressure and infections that go unnoticed due to a lack of medical facilities and trained professionals. Early marriages, limited education, and poverty add to the problem. Women often cannot reach healthcare centers in time, leading to complications which increases the risk associated with the same. This project develops machine learning models to predict maternal health risk using vital health indicators, enabling early identification of high-risk pregnancies.

Contents

1	Introduction	2
1.1	Problem Statement	2
1.2	Dataset Description	2
2	Exploratory Data Analysis	2
2.1	Data Quality and Preprocessing	2
2.2	Summary Statistics	4
3	Mathematical Overview	5
3.1	Random Forest	5
3.2	Support Vector Machine (SVM)	5
4	Model Fitting and Comparison	5
4.1	Data Splitting and Cross-Validation	5
4.2	Handling Class Imbalance	6
4.3	Model Training	6
4.4	Test Set Evaluation	7
5	Interpretable Machine Learning (XAI)	9
5.1	Feature Importance	9
5.2	Partial Dependence Plots	9
5.3	Local Explanations (LIME)	11
6	Conclusions	11
6.1	Summary	11
6.2	Model Comparison Results	11
6.3	Key Findings	12
6.4	Clinical Recommendations	12
6.5	Limitations	12
7	References	12

1 Introduction

1.1 Problem Statement

Maternal mortality remains a critical global health challenge. This project develops machine learning models to predict maternal health risk as a **binary classification** (High Risk vs. Not High Risk) based on vital health indicators.

Rationale for Binary Classification: The original dataset contains three ordinal risk levels (low, mid, high). Since ordinal relationships are not optimally captured by standard multi-class classifiers, we aggregate mid and low risk into “Not High Risk”. This directly addresses: *“Is this pregnancy high-risk?”*

1.2 Dataset Description

The Maternal Health Risk dataset was collected from hospitals in rural Bangladesh via an IoT-based monitoring system (Ahmed and Kashem 2023). It contains 1,014 observations with 6 predictor variables.

Variable	Description	Range
Age	Age of pregnant woman (years)	10-70
SystolicBP	Systolic blood pressure (mmHg)	70-160
DiastolicBP	Diastolic blood pressure (mmHg)	49-100
BS	Blood sugar level (mmol/L)	6.0-19.0
BodyTemp	Body temperature (°F)	98-103
HeartRate	Heart rate (bpm)	7-90
RiskLevel	Target: High Risk vs. Not High Risk	2 classes

2 Exploratory Data Analysis

2.1 Data Quality and Preprocessing

The dataset has **no missing values**. Two observations with HeartRate = 7 bpm (physiologically impossible) were removed, leaving **1012 observations**.

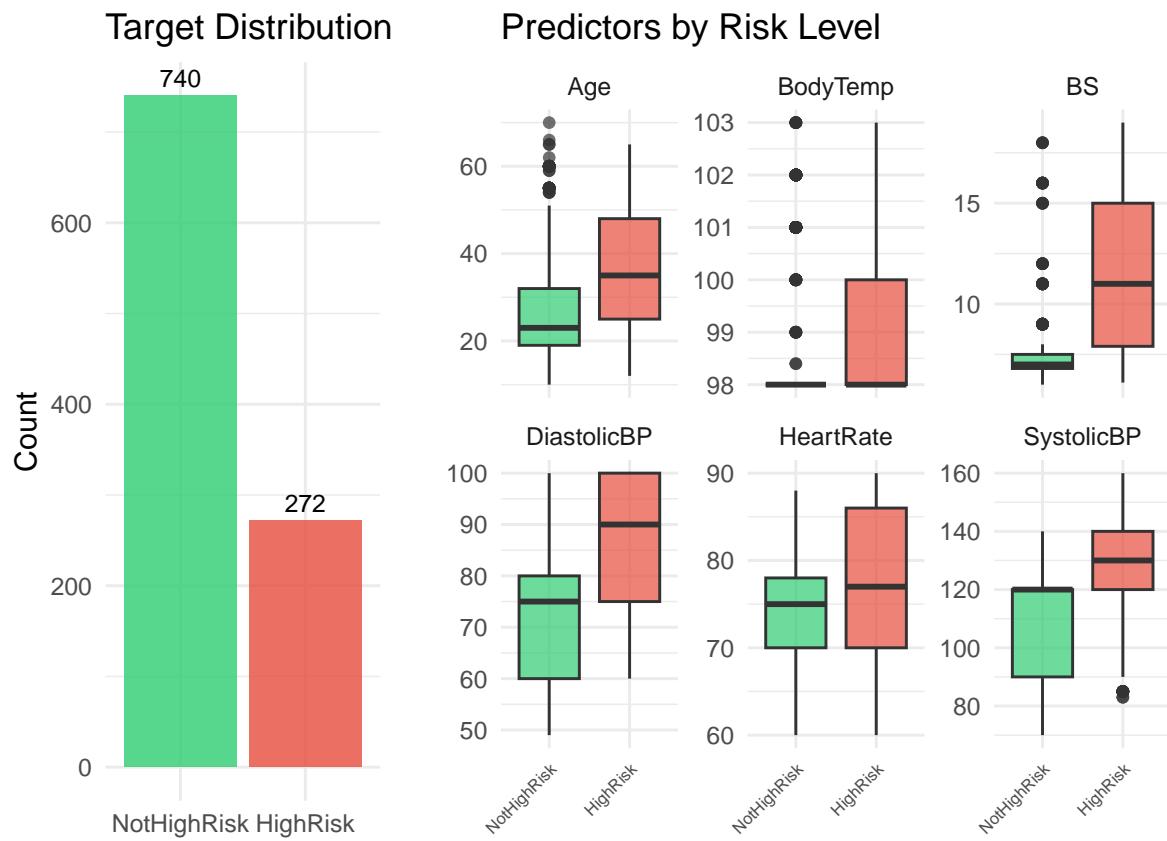


Figure 1: Target Distribution and Predictor Variables by Risk Level

Key Observations: Blood Sugar (BS) and Systolic BP are strong discriminators for high-risk cases. Class imbalance (~27% HighRisk) is addressed using stratified sampling.

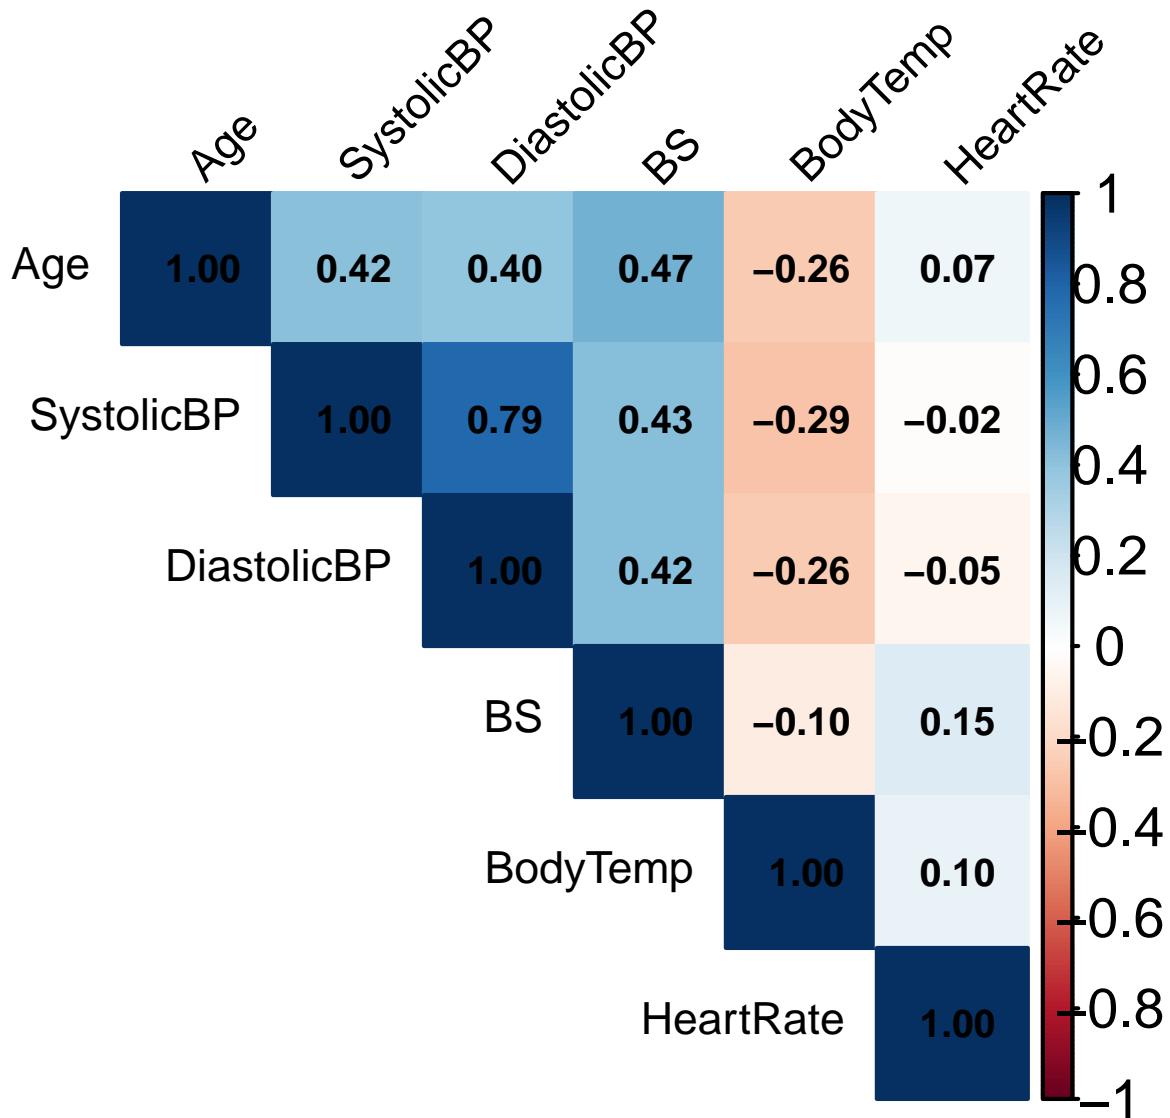


Figure 2: Correlation Matrix

No severe multicollinearity detected. Systolic and Diastolic BP show expected moderate correlation.

2.2 Summary Statistics

Table 2: Summary Statistics of Predictor Variables

Variable	Min	Mean	Max	SD
Age	10	29.9	70	13.5
SystolicBP	70	113.2	160	18.4
DiastolicBP	49	76.5	100	13.9
BS	6	8.7	19	3.3
BodyTemp	98	98.7	103	1.4
HeartRate	60	74.4	90	7.5

3 Mathematical Overview

3.1 Random Forest

Random Forest is an ensemble learning method that constructs multiple decision trees during training (Breiman 2001). The algorithm works as follows:

1. **Bootstrap Sampling:** For each tree, draw a bootstrap sample (with replacement) from the training data
2. **Feature Randomization:** At each node split, only consider a random subset of $m = \sqrt{p}$ features
3. **Tree Construction:** Build each tree to maximum depth without pruning
4. **Aggregation:** Combine predictions via majority voting (classification)

Prediction Formula:

$$\hat{f}(x) = \text{mode}\{h_1(x), h_2(x), \dots, h_B(x)\}$$

where $h_b(x)$ is the prediction of tree b and B is the total number of trees.

Split Criterion - Gini Impurity:

$$G(t) = 1 - \sum_{k=1}^K p_k^2$$

where p_k is the proportion of class k observations at node t . A split is chosen to maximize the reduction in impurity.

Key Hyperparameters: `ntree` (number of trees), `mtry` (features per split), `nodesize` (minimum node size).

3.2 Support Vector Machine (SVM)

SVM finds the optimal separating hyperplane that maximizes the margin between classes (James et al. 2021). For non-linearly separable data, the soft-margin SVM solves:

Optimization Problem:

$$\min_{\beta, \beta_0} \frac{1}{2} \|\beta\|^2 + C \sum_{i=1}^n \xi_i$$

subject to: $y_i(\beta^T x_i + \beta_0) \geq 1 - \xi_i$ and $\xi_i \geq 0$

where C controls the trade-off between margin maximization and misclassification penalty, and ξ_i are slack variables.

RBF (Radial Basis Function) Kernel:

$$K(x_i, x_j) = \exp(-\gamma \|x_i - x_j\|^2)$$

The RBF kernel maps data to infinite-dimensional space, enabling non-linear decision boundaries. Parameter γ controls the kernel width: larger γ means tighter fit (risk of overfitting).

Key Hyperparameters: C (cost/regularization), γ (kernel width, often denoted as `sigma` in R).

4 Model Fitting and Comparison

4.1 Data Splitting and Cross-Validation

```
## Original Training Data Class Distribution:  
##  
## NotHighRisk    HighRisk  
##      592        218
```

Data Split Strategy: Following professor's guidelines, since we use 10-fold cross-validation for hyperparameter tuning, we combine training and validation sets ($60\% + 20\% = 80\%$). The test set (20%) is held out exclusively for final model comparison.

4.2 Handling Class Imbalance

The dataset shows class imbalance (~27% HighRisk vs ~73% NotHighRisk). To prevent bias toward the majority class, we apply **oversampling** to balance the training data.

```
## Balanced Training Data Class Distribution:  
##  
## NotHighRisk    HighRisk  
##      592        592
```

After oversampling, the training data is balanced with equal representation of both classes. This ensures the models learn to identify high-risk cases effectively without being biased toward the majority class. The test set remains **unbalanced** to reflect real-world class distribution for fair evaluation.

4.3 Model Training

We trained three models (Random Forest, Decision Tree, SVM) and selected the **best 2 based on CV AUC-ROC** for final comparison.

```
## Model Selection (CV AUC-ROC):  
## 1. Random Forest : 0.9856  
## 2. SVM : 0.9777  
## 3. Decision Tree : 0.9614  
##  
## Best 2 models selected: Random Forest and SVM
```

Why Random Forest and SVM?

We selected these two models based on multiple criteria:

1. **Cross-Validation Performance:** Random Forest achieved the highest CV AUC-ROC, followed by SVM. Decision Tree showed lower performance, likely due to its tendency to overfit without ensemble averaging.
2. **Algorithmic Diversity:** RF and SVM represent fundamentally different approaches:
 - **Random Forest:** Ensemble of decision trees using bagging and feature randomization. Handles non-linear relationships naturally and provides built-in feature importance.
 - **SVM:** Finds optimal separating hyperplane in transformed feature space using the RBF kernel. Effective for smaller datasets and robust to outliers due to margin maximization.
3. **Complementary Strengths:** RF excels at capturing complex feature interactions, while SVM with RBF kernel provides smooth decision boundaries. Comparing both allows us to assess whether findings are model-specific or generalizable.
4. **Clinical Applicability:** Both models support probability outputs needed for risk stratification, and both can be interpreted using model-agnostic XAI techniques (feature importance, PDPs, LIME).

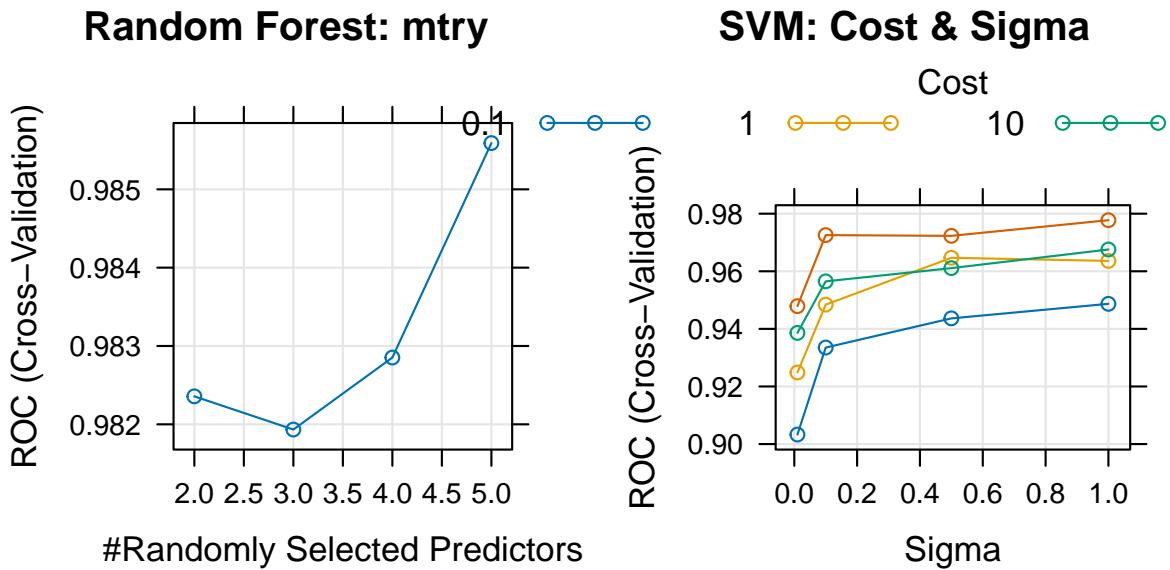


Figure 3: Hyperparameter Tuning Results (Best 2 Models)

Best Hyperparameters: RF: mtry = 5; SVM: C = 100, sigma = 1

4.4 Test Set Evaluation

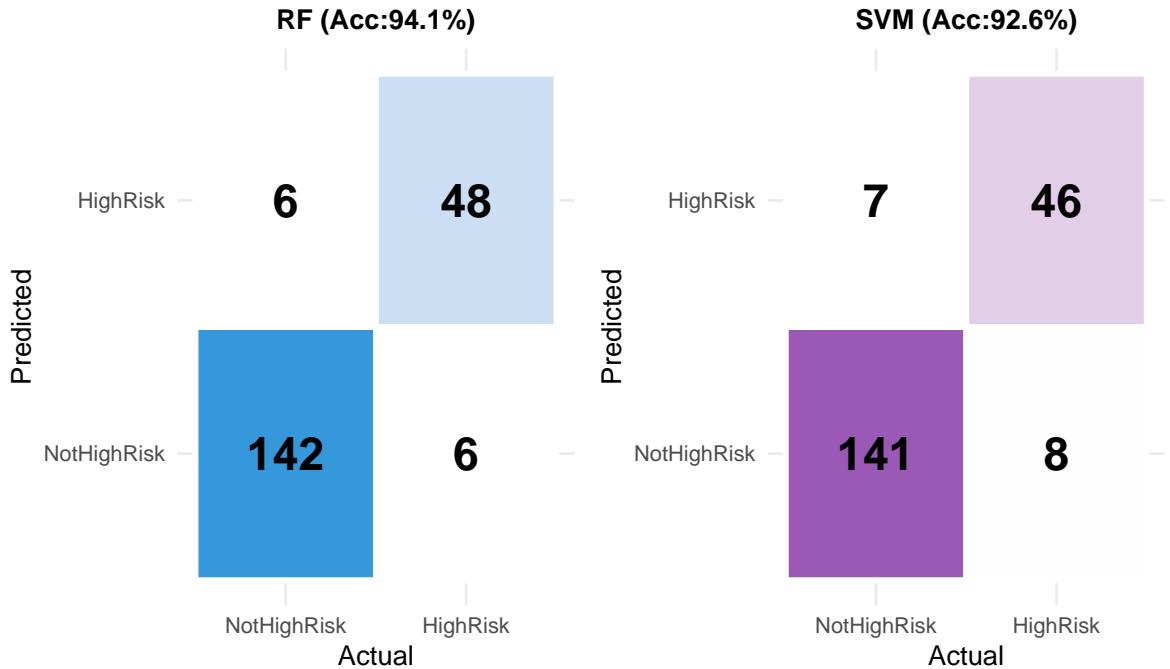


Figure 4: Confusion Matrices: Random Forest (left) and SVM (right)

Table 3: Test Set Performance Metrics (in percentage)

	Metric	RF	SVM
Accuracy	Accuracy	94.1	92.6
Sensitivity	Sensitivity	95.9	95.3
Specificity	Specificity	88.9	85.2
Pos Pred Value	Precision	95.9	94.6

	Metric	RF	SVM
F1	F1 Score	95.9	94.9

ROC Curves

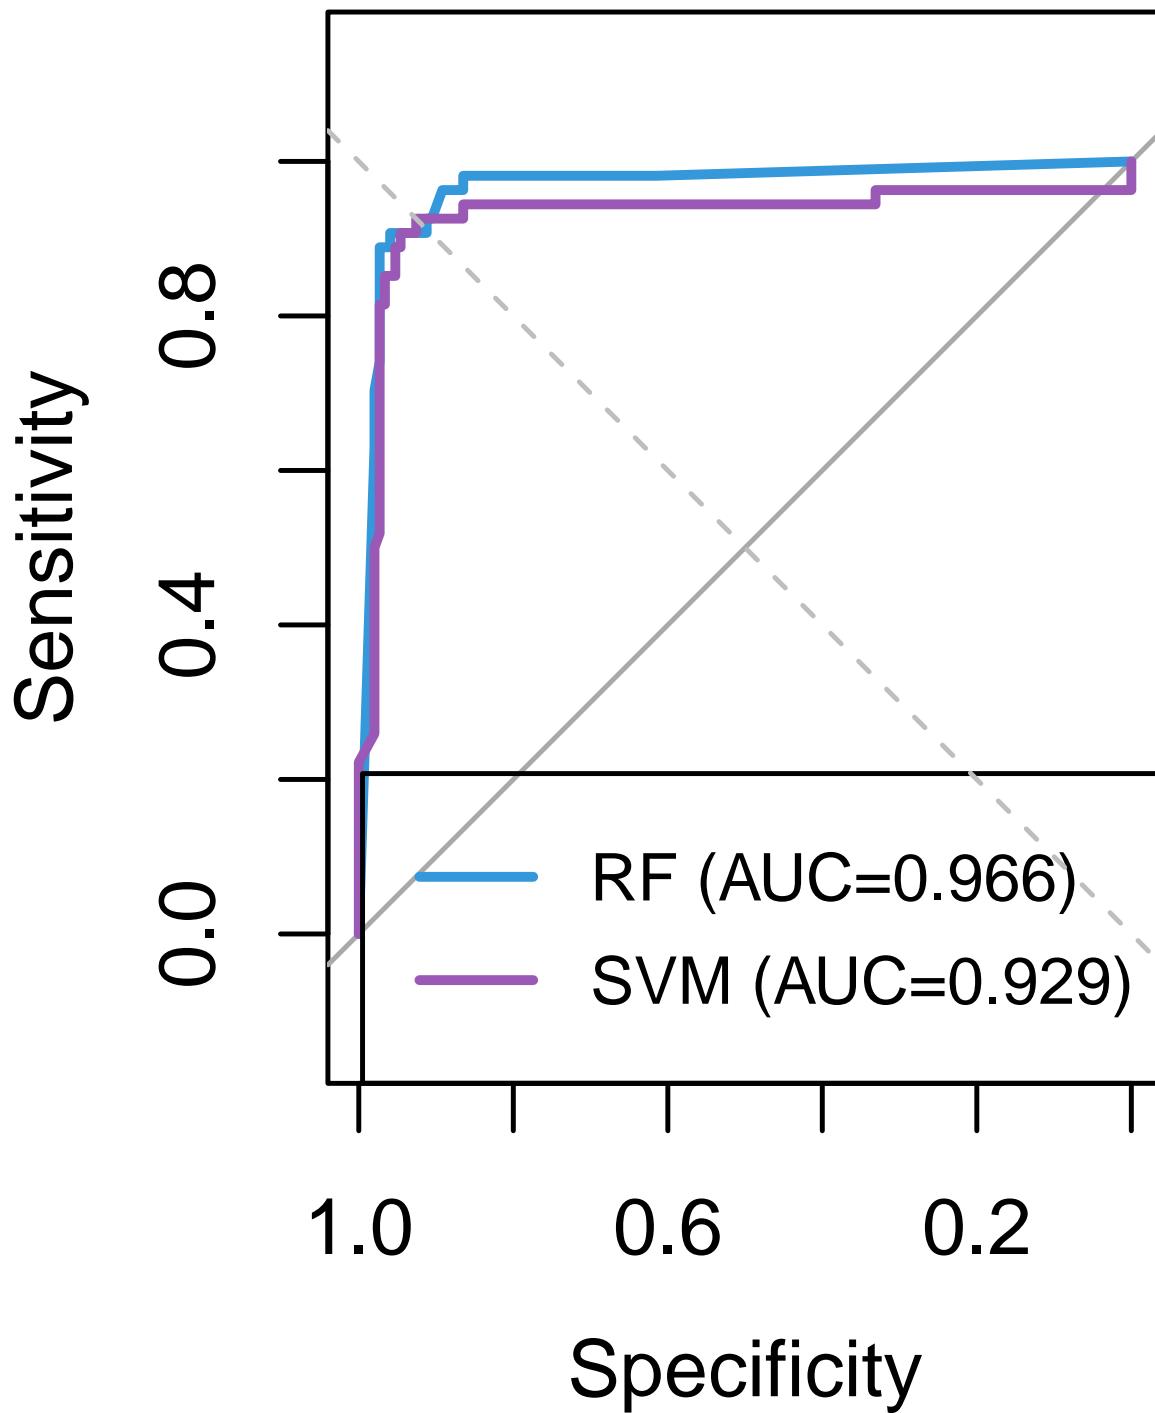


Figure 5: ROC Curves and Performance Comparison

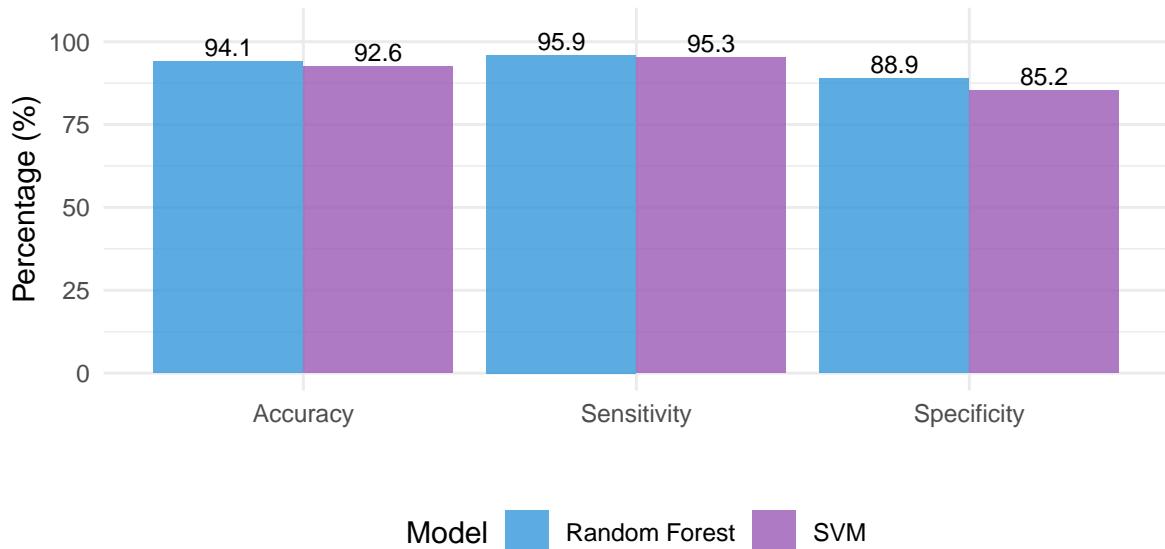


Figure 6: Model Performance Comparison

5 Interpretable Machine Learning (XAI)

5.1 Feature Importance

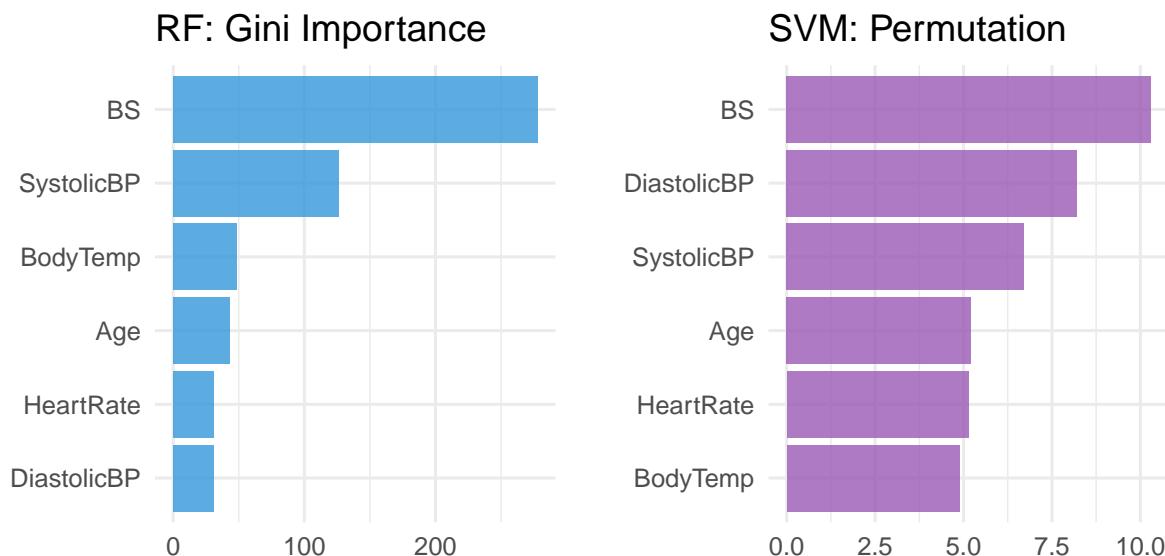


Figure 7: Feature Importance Comparison

Both models rank **Blood Sugar (BS)** as the most important feature, followed by **SystolicBP** and **Age**.

5.2 Partial Dependence Plots

Partial Dependence Plots (PDPs) show the marginal effect of a feature on the predicted outcome, averaging over all other features. We compare PDPs for both Random Forest and SVM models.

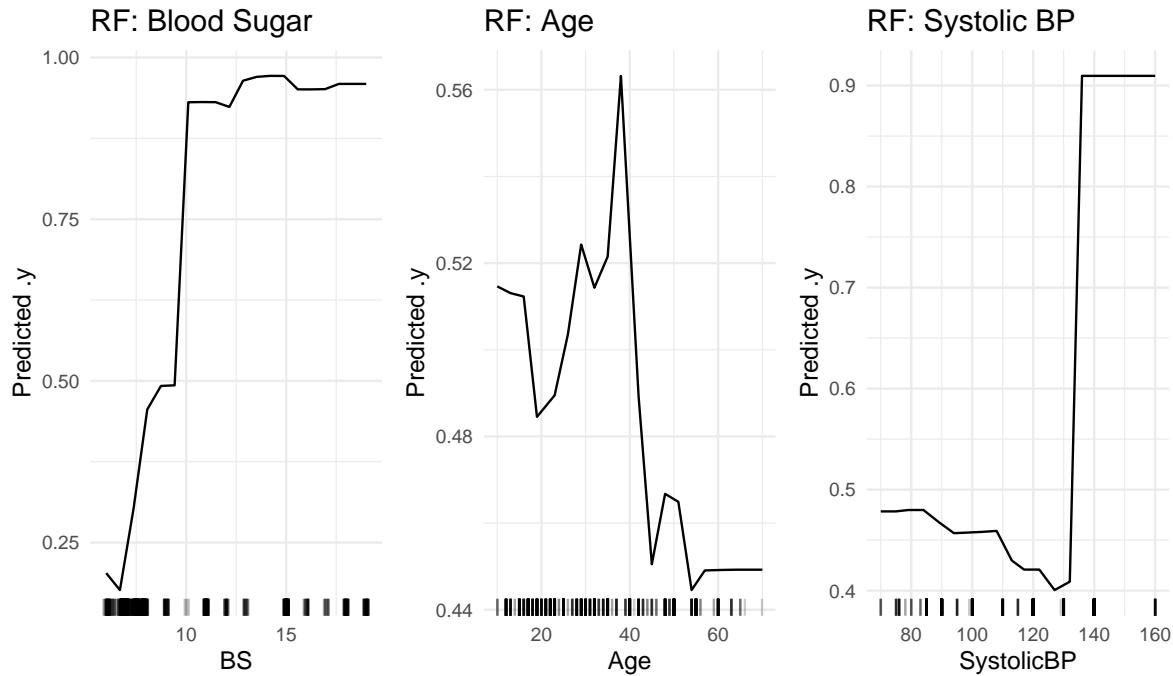


Figure 8: Random Forest: Partial Dependence Plots for Top 3 Features

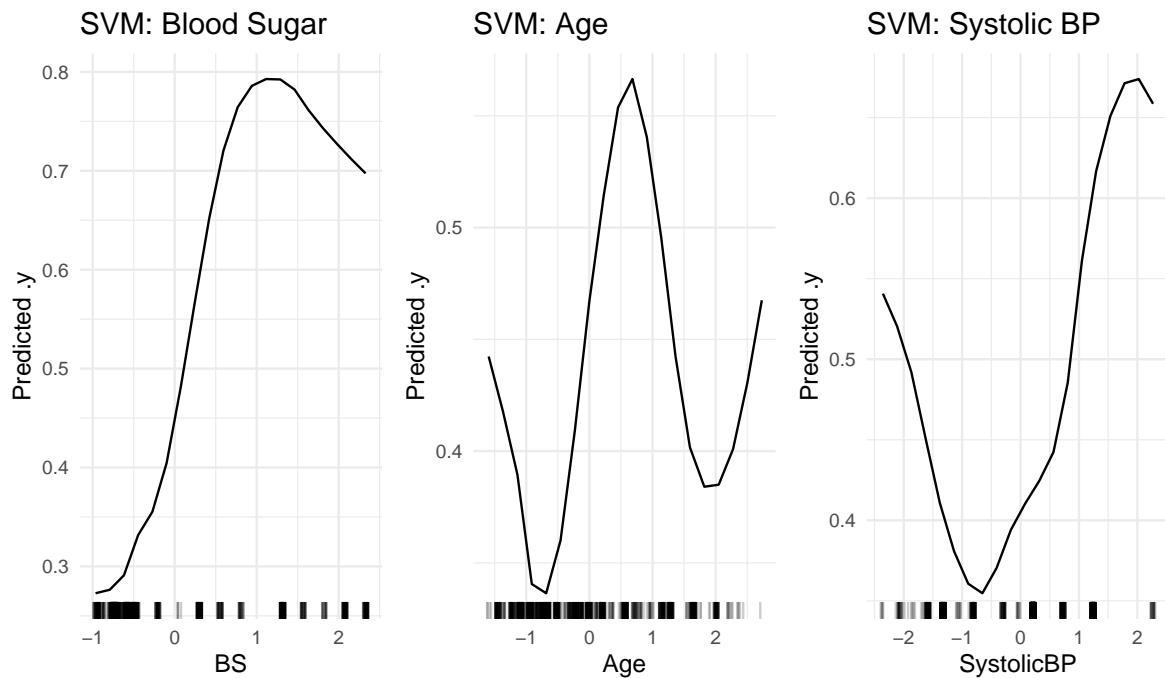


Figure 9: SVM: Partial Dependence Plots for Top 3 Features

PDP Comparison:

- **Blood Sugar (BS):** Both models show strong positive relationship. RF shows step-like pattern (tree-based), while SVM shows smoother transitions due to the RBF kernel.
- **Age:** Both models show moderate positive effect, with RF displaying more abrupt changes and SVM showing gradual increase.
- **Systolic BP:** Both identify ~130 mmHg as a risk threshold, but SVM captures a more continuous relationship.

5.3 Local Explanations (LIME)

LIME (Local Interpretable Model-agnostic Explanations) provides instance-level explanations by fitting a simple interpretable model locally around a prediction. We compare LIME explanations for both models on the same high-risk case.

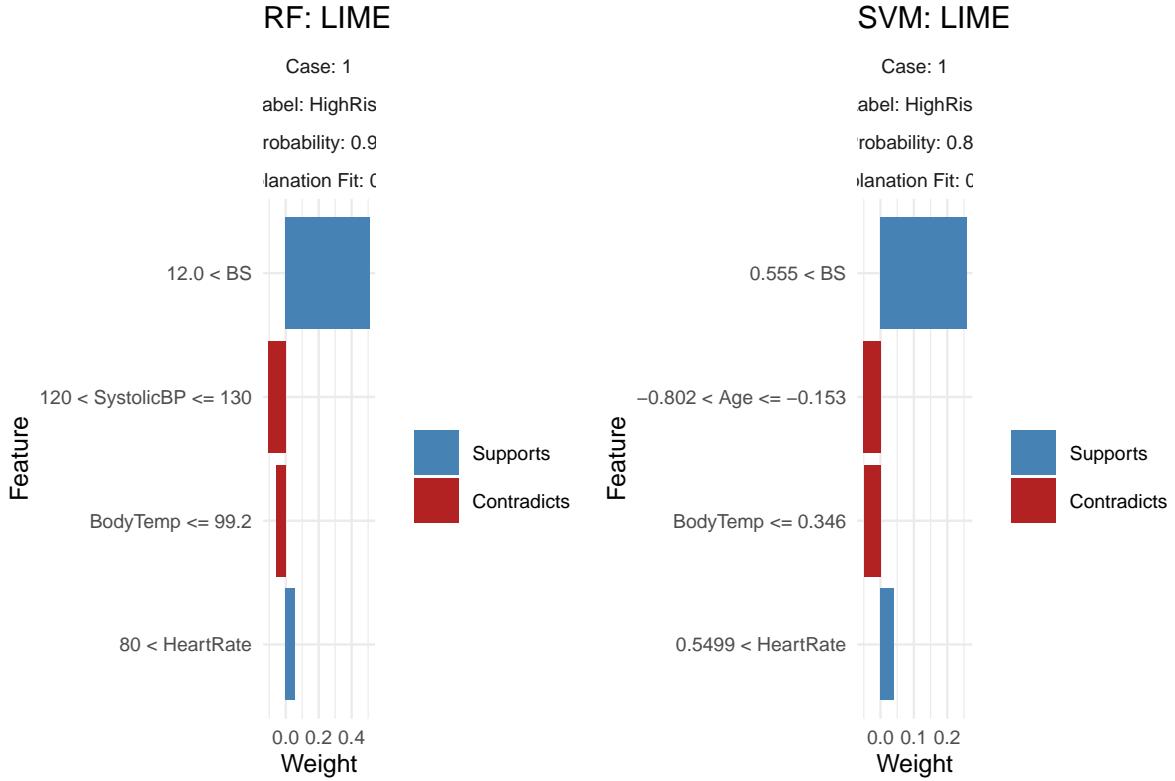


Figure 10: LIME Explanations: RF (left) vs SVM (right) for Same High-Risk Case

LIME Comparison: Both models identify similar key features for this high-risk case, but with different contribution magnitudes. This consistency across models strengthens confidence in the clinical interpretation.

6 Conclusions

6.1 Summary

We trained three machine learning models (Random Forest, Decision Tree, SVM) and selected the **best 2 (Random Forest and SVM)** based on cross-validation AUC-ROC performance. Both selected models achieve strong performance in detecting high-risk pregnancies using binary classification.

6.2 Model Comparison Results

Table 4: Final Model Comparison Summary

Metric	RF	SVM
CV AUC-ROC	0.986	0.978
Test AUC-ROC	0.966	0.929
Accuracy	94.1%	92.6%
Sensitivity	95.9%	95.3%
Specificity	88.9%	85.2%
F1 Score	95.9%	94.9%

6.3 Key Findings

1. **Random Forest outperforms SVM** across most metrics, particularly in sensitivity which is critical for detecting high-risk cases
2. **Blood Sugar (BS) is the most important predictor** across both models, consistent with medical literature on gestational diabetes
3. **Systolic Blood Pressure and Age** are secondary important features, aligning with known risk factors for pregnancy complications
4. **Both models achieve excellent discrimination** with AUC-ROC > 0.90, indicating reliable separation between risk classes

6.4 Clinical Recommendations

Based on our analysis, we recommend:

- **Primary Screening:** Use Random Forest model for initial risk assessment due to higher sensitivity
- **Key Indicators to Monitor:** Blood sugar levels should be closely monitored, especially values > 8 mmol/L
- **Blood Pressure Monitoring:** Systolic BP > 130 mmHg should trigger additional evaluation
- **Age Consideration:** Older maternal age warrants closer monitoring

6.5 Limitations

- Dataset size (~1,000 observations) may limit generalizability
- Geographic scope limited to rural Bangladesh
- Limited feature set (6 predictors) - additional clinical variables could improve predictions
- Binary classification loses granularity of original ordinal risk levels

7 References

Ahmed, Marzia, and Mohammod Abul Kashem. 2023. *Maternal Health Risk Data Set*. UCI Machine Learning Repository. <https://archive.ics.uci.edu/dataset/863/maternal+health+risk>.

Breiman, Leo. 2001. “Random Forests.” *Machine Learning* 45 (1): 5–32. <https://doi.org/10.1023/A:1010933404324>.

James, Gareth, Daniela Witten, Trevor Hastie, and Robert Tibshirani. 2021. *An Introduction to Statistical Learning: With Applications in r*. 2nd ed. Springer.